


16-BIT ML

K-Means Clustering

How feasible is basic clustering with 1000's of points on vintage hardware?



ALGORITHM

- Make initial guess at cluster centroids
- Compute distance from every point to every centroid
- Assign points to closest centroid
- Compute new centroids
- Repeat

DEMO

- Display shows a non-uniform random 2D distribution
- Color indicates current cluster membership ($K=6$)
- White crosses show current cluster centroids
- Watch for changes from frame to frame



IMPLEMENTATION

- Turbo Pascal 5.0 on IBM PCjr (~1984)
- CPU: 8088 @ 4.77 MHz
- RAM: 64+64+640 kB

